

CS 3200

Database Design

Nik Bear Brown

nik@ccs.neu.edu

SQL

Topics

- Relational Algebra
- Relational Calculus

Relational algebra and relational calculus

- Relational algebra and relational calculus are formal languages associated with the relational model.
- Informally, relational algebra is a (high-level) procedural language and relational calculus a non-procedural language.
- However, formally both are equivalent to one another.
- A language that produces a relation that can be derived using relational calculus is relationally complete.

SQL Overview

- SQL is the language of relational databases.
- It is used for every aspect of database development and management.
- Anyone who works with relational databases is expected to have a knowledge of SQL.

Objectives of SQL

- Ideally, database language should allow user to:
 - create the database and relation structures;
 - perform insertion, modification, deletion of data from relations;
 - perform simple and complex queries.
- Must perform these tasks with minimal user effort and command structure/syntax must be easy to learn.
- It must be portable.

Non-procedural SQL

It is non-procedural - you specify *what* information you require, rather than *how* to get it;

Objectives of SQL

- Consists of standard English words:

```
1) CREATE TABLE Staff(staffNo VARCHAR(5),  
                        IName VARCHAR(15),  
                        salary DECIMAL(7,2));
```

```
2) INSERT INTO Staff VALUES ('SG16', 'Brown', 8300);
```

```
3) SELECT staffNo, IName, salary  
   FROM Staff  
   WHERE salary > 10000;
```

History of SQL

- In 1974, D. Chamberlin (IBM San Jose Laboratory) defined language called 'Structured English Query Language' (SEQUEL).
- A revised version, SEQUEL/2, was defined in 1976 but name was subsequently changed to SQL for legal reasons.

History of SQL

- Still pronounced 'see-quel', though official pronunciation is 'S-Q-L'. (or squirrel or 's-quel',)
- IBM subsequently produced a prototype DBMS called *System R*, based on SEQUEL/2.
- Roots of SQL, however, are in SQUARE (Specifying Queries as Relational Expressions), which predates System R project.

History of SQL

- In late 70s, ORACLE appeared and was probably first commercial RDBMS based on SQL.
- In 1987, ANSI and ISO published an initial standard for SQL.
- In 1989, ISO published an addendum that defined an 'Integrity Enhancement Feature'.
- In 1992, first major revision to ISO standard occurred, referred to as SQL2 or SQL/92.
- In 1999, SQL:1999 was released with support for object-oriented data management.
- In late 2003, SQL:2003 was released.
- In summer 2008, SQL:2008 was released.
- In late 2011, SQL:2011 was released.

SQL History

- SQL is the programming language used for accessing and manipulating data and objects in relational databases.
- The first versions of SQL were developed by IBM in the 1970s.
- SQL first became an ANSI standard in 1986 and an ISO standard in 1987.
- There was a major revision to the standard in 1992.
- Additional modifications were made in 1999, 2003, and 2006.

Importance of SQL

- SQL has become part of application architectures such as IBM's Systems Application Architecture.
- It is strategic choice of many large and influential organizations (e.g. X/OPEN).
- SQL is Federal Information Processing Standard (FIPS) to which conformance is required for all sales of databases to American Government.

Nature of SQL

- SQL is a declarative language.
- Procedural languages like C# or Java describe how to accomplish a task step by step.
- In a declarative language, you say *what* you want to do, not *how*.

SQL Functionality

- SQL is not case sensitive.
- In some environments SQL statements must be ended with a semicolon.
- SQL is usually divided into two broad areas of functionality:
 - DDL (Data Definition Language)
 - DML (Data Manipulation Language)

DDL

- Data definition language is the set of SQL keywords and commands used to create, alter, and remove database objects.
- An example is the `CREATE TABLE` command:

```
CREATE TABLE TestTable  
(  
    TestID INT IDENTITY (1,1),  
    TestDescription NVARCHAR(255)  
)
```

DML

- Data manipulation language is the set of key words and commands used to retrieve and modify data.
- `SELECT`, `UPDATE`, `INSERT`, and `DELETE` are the primary actions of DML.

Importance of SQL

- SQL is used in other standards and even influences development of other standards as a definitional tool. Examples include:

- ISO's Information Resource Directory System (IRDS) Standard
- Remote Data Access (RDA) Standard.

Writing SQL Commands

- SQL statement consists of *reserved words* and *user-defined words*.
 - Reserved words are a fixed part of SQL and must be spelt exactly as required and cannot be split across lines.
 - User-defined words are made up by user and represent names of various database objects such as relations, columns, views.

Writing SQL Commands

- Most components of an SQL statement are *case insensitive*, except for literal character data.
- More readable with indentation and lineation:
 - Each clause should begin on a new line.
 - Start of a clause should line up with start of other clauses.
 - If clause has several parts, should each appear on a separate line and be indented under start of clause.

Writing SQL Commands

- Use extended form of BNF notation:
 - Upper-case letters represent reserved words.
 - Lower-case letters represent user-defined words.
 - | indicates a *choice* among alternatives.
 - Curly braces indicate a *required element*.
 - Square brackets indicate an *optional element*.
 - ... indicates *optional repetition* (0 or more).

Literals

- Literals are constants used in SQL statements.
- All non-numeric literals must be enclosed in single quotes (e.g. 'London').
- All numeric literals must not be enclosed in quotes (e.g. 650.00).

Select Statement

- The SELECT statement is used to retrieve data from the database.
- The basic syntax is:

```
SELECT <columnName>, <columnName>  
FROM <TableName>
```

```
SELECT StudentFirstName, StudentLastName, StudentPhone  
FROM Student
```

SELECT Statement

```
SELECT [DISTINCT | ALL]
      { * | [columnExpression [AS newName]] [, ...]
      }
FROM      TableName [alias] [, ...]
[WHERE    condition]
[GROUP BY columnList] [HAVING condition]
[ORDER BY columnList]
```

SELECT Statement

FROM Specifies table(s) to be used.

WHERE Filters rows.

GROUP BY Forms groups of rows with same column value.

HAVING Filters groups subject to some condition.

SELECT Specifies which columns are to appear in output.

ORDER BY Specifies the order of the output.

SELECT Statement

- Order of the clauses cannot be changed.
- Only SELECT and FROM are mandatory.

Select All Columns, All Rows

List full details of all staff.

```
SELECT    staffNo,    fName,    lName,  
address,  
    position, sex, DOB, salary, branchNo  
FROM Staff;
```

- Can use * as an abbreviation for 'all columns':

```
SELECT *  
FROM Staff;
```

Select Comparison Search Condition

List all staff with a salary greater than 10,000.

```
SELECT staffNo, fName, lName, position,  
salary  
FROM Staff  
WHERE salary > 10000;
```

staffNo	fName	lName	position	salary
SL21	John	White	Manager	30000.00
SG37	Ann	Beech	Assistant	12000.00
SG14	David	Ford	Supervisor	18000.00
SG5	Susan	Brand	Manager	24000.00

Select Compound Comparison Search Condition

List addresses of all branch offices in London or Glasgow.

```
SELECT *  
FROM Branch  
WHERE city = 'London' OR city = 'Glasgow';
```

branchNo	street	city	postcode
B005	22 Deer Rd	London	SW1 4EH
B003	163 Main St	Glasgow	G11 9QX
B002	56 Clover Dr	London	NW10 6EU

Select Range Search Condition

List all staff with a salary between 20,000 and 30,000.

```
SELECT    staffNo,    fName,    lName,  
position, salary  
FROM Staff  
WHERE salary BETWEEN 20000 AND 30000;
```

- BETWEEN test includes the endpoints of range.

Select Range Search Condition

- Also a negated version NOT BETWEEN.
- BETWEEN does not add much to SQL's expressive power. Could also write:

```
SELECT staffNo, fName, lName, position, salary
FROM Staff
WHERE salary >= 20000 AND salary <= 30000;
```

- Useful, though, for a range of values.

Select Set Membership

List all managers and supervisors.

```
SELECT staffNo, fName, lName, position
FROM Staff
WHERE position IN ('Manager',
                  'Supervisor');
```

staffNo	fName	lName	position
SL21	John	White	Manager
SG14	David	Ford	Supervisor
SG5	Susan	Brand	Manager

The * WildCard

- Instead of listing each of the columns, you can use an * to include all columns.

`SELECT * FROM`

- Listing the columns does give you the ability to choose both which columns and which order to present them.
- With the * you return all the columns in the order they have in the underlying table.

Distinct Key Word

- Sometimes a query will return multiple duplicate values.
- For instance the statement

```
SELECT Key  
FROM Session
```

Could return numerous instances of each customer.
- The **DISTINCT** keyword will make it so it only returns one instance of each Key.

Distinct Key Word

```
SELECT DISTINCT Key  
FROM Session
```

- The `DISTINCT` keyword always operates on the whole row, not on individual columns.
- It only returns distinct rows.

Calculations

- You can do calculations in SELECT statements.

```
SELECT ItemNumber, ItemPrice, Quantity, ItemPrice *  
Quantity  
FROM CustomerOrder
```

Order of Operations

The order of operation is the same as in algebra.

1. Whatever is in parentheses is executed first. If parentheses are nested, the innermost is executed first, then the next most inner, etc.
2. Then all division and multiplication left to right
3. And finally all addition and subtraction left to right

Sorting

- You can sort the results of a query by using the keywords ORDER BY.
`SELECT *`
`FROM Session`
`ORDER BY SessionDate`
- ORDER BY does an ascending A-Z, 1-10, etc. sort by default.
- You can change the direction by using the DESC keyword after the field to be sorted.

Aliasing

- Sometimes it is useful to alias a column name to make a more readable result set.

```
SELECT StudentLastName AS [Last Name], StudentFirstName AS  
[First Name]  
FROM Student
```

- The `AS` keyword is optional.
- Double quotes “ ” can be used instead of square brackets.

Where Clause

- The `WHERE` clause allows you to limit the rows you return in a query.
- You use the `WHERE` clause to specify the criteria by which the rows will be filtered.

```
SELECT LastName, FirstName, Phone, City  
FROM Customer  
WHERE City = 'Boston'
```

Other Criteria

- As well as equal you can use other operators for the criteria:
 - >
 - <
 - >=
 - =<
- Character and date values in the criteria are quoted with single quotes.
- Numerical values are not quoted.

Like

- The `LIKE` keyword used in a `WHERE` operator with a wildcard (`%` or `_`) allows you to search for patterns in character-based fields.
- The following returns all items whose name starts with “T.”

```
SELECT ItemName, ItemPrice
FROM Inventory
WHERE ItemName LIKE 'T%'
```

Between

- The `BETWEEN` keyword can be used in criteria to return values between to other values.
- `BETWEEN` is inclusive of its ends.

```
SELECT Key, SessionDate, StudentKey  
FROM Session  
WHERE SessionDate BETWEEN '11/1/2014' AND  
      '11/30/2014'
```

AND OR NOT

- You can use keywords `AND`, `OR`, and `NOT` to combine criteria in a query.
- `AND` is exclusive. `Or` is Inclusive.
- `WHERE City = 'Boston' OR City='Los Angeles'` returns all records that have either Boston or Los Angeles for their city.
- `WHERE City='Boston' AND City='Los Angeles'` returns nothing because the record cannot have both at the same time.
- `NOT` excludes.
- `WHERE NOT City = 'Los Angeles'` returns every city except Los Angeles.

NULL

- Nulls are special cases. They are not a value and so cannot be compared to a value using = or < or >.
- To locate nulls you can use the IS keyword in a criteria:

WHERE StudentKey IS NULL

WHERE StudentKey IS NOT NULL

Select Set Membership

- There is a negated version (NOT IN).
- IN does not add much to SQL's expressive power. Could have expressed this as:

```
SELECT staffNo, fName, lName, position  
FROM Staff  
WHERE position='Manager' OR  
       position='Supervisor';
```

- IN is more efficient when set contains many values.

Select Pattern Matching

Find all owners with the string 'Glasgow' in their address.

```
SELECT ownerNo, fName, lName, address,  
telNo  
FROM PrivateOwner  
WHERE address LIKE '%Glasgow%';
```

ownerNo	fName	lName	address	telNo
CO87	Carol	Farrel	6 Achray St, Glasgow G32 9DX	0141-357-7419
CO40	Tina	Murphy	63 Well St, Glasgow G42	0141-943-1728
CO93	Tony	Shaw	12 Park Pl, Glasgow G4 0QR	0141-225-7025

Select Pattern Matching

- SQL has two special pattern matching symbols:
 - %: sequence of zero or more characters;
 - _ (underscore): any single character.
- LIKE '%Glasgow%' means a sequence of characters of any length containing '*Glasgow*'.

Select NULL Search Condition

List details of all viewings on property PG4 where a comment has not been supplied.

- There are 2 viewings for property PG4, one with and one without a comment.
- Have to test for null explicitly using special keyword IS NULL:

```
SELECT clientNo, viewDate
FROM Viewing
WHERE propertyNo = 'PG4' AND
               comment IS NULL;
```


Select NULL Search Condition

clientNo	viewDate
CR56	26-May-04

- Negated version (IS NOT NULL) can test for non-null values.

Select Single Column Ordering

List salaries for all staff, arranged in descending order of salary.

```
SELECT staffNo, fName, lName, salary
FROM Staff
ORDER BY salary DESC;
```

Select Multiple Column Ordering

Produce abbreviated list of properties in order of property type.

```
SELECT propertyNo, type, rooms, rent  
FROM PropertyForRent  
ORDER BY type;
```

Select Multiple Column Ordering

To arrange in order of rent, specify minor order:

```
SELECT propertyNo, type, rooms, rent  
FROM PropertyForRent  
ORDER BY type, rent DESC;
```

SELECT Statement - Aggregates

- ISO standard defines five aggregate functions:

COUNT returns number of values in specified column.

SUM returns sum of values in specified column.

AVG returns average of values in specified column.

MIN returns smallest value in specified column.

MAX returns largest value in specified column.

SELECT Statement - Aggregates

- Each operates on a single column of a table and returns a single value.
- COUNT, MIN, and MAX apply to numeric and non-numeric fields, but SUM and AVG may be used on numeric fields only.
- Apart from COUNT(*), each function eliminates nulls first and operates only on remaining non-null values.

SELECT Statement - Aggregates

- COUNT(*) counts all rows of a table, regardless of whether nulls or duplicate values occur.
- Can use DISTINCT before column name to eliminate duplicates.
- DISTINCT has no effect with MIN/MAX, but may have with SUM/AVG.

SELECT Statement - Aggregates

- Aggregate functions can be used only in SELECT list and in HAVING clause.
- If SELECT list includes an aggregate function and there is no GROUP BY clause, SELECT list cannot reference a column out with an aggregate function. For example, the following is illegal:

```
SELECT staffNo, COUNT(salary)
FROM Staff;
```


Select - Use of COUNT(*)

How many properties cost more than £350 per month to rent?

```
SELECT COUNT(*) AS myCount  
FROM PropertyForRent  
WHERE rent > 350;
```

myCount
5

Select - Use of COUNT(DISTINCT)

How many different properties viewed in May '13?

```
SELECT COUNT(DISTINCT propertyNo) AS  
myCount  
FROM Viewing  
WHERE viewDate BETWEEN '1-May-13'  
AND '31-May-13';
```

myCount
2

Select - Use of COUNT and SUM

Find number of Managers and sum of their salaries.

```
SELECT COUNT(staffNo) AS myCount,  
       SUM(salary) AS mySum  
FROM Staff  
WHERE position = 'Manager';
```

myCount	mySum
2	54000.00

Select - Use of MIN, MAX, AVG

Find minimum, maximum, and average staff salary.

```
SELECT      MIN(salary)      AS  
myMin,  
            MAX(salary) AS myMax,  
            AVG(salary) AS myAvg  
FROM Staff;
```

myMin	myMax	myAvg
9000.00	30000.00	17000.00

SELECT Statement - Grouping

- Use GROUP BY clause to get sub-totals.
- SELECT and GROUP BY closely integrated: each item in SELECT list must be *single-valued per group*, and SELECT clause may only contain:
 - column names
 - aggregate functions
 - constants
 - expression involving combinations of the above.

SELECT Statement - Grouping

- All column names in SELECT list must appear in GROUP BY clause unless name is used only in an aggregate function.
- If WHERE is used with GROUP BY, WHERE is applied first, then groups are formed from remaining rows satisfying predicate.
- ISO considers two nulls to be equal for purposes of GROUP BY.

Select - Use of GROUP BY

Find number of staff in each branch and their total salaries.

```
SELECT branchNo,  
        COUNT(staffNo)      AS  
myCount,  
        SUM(salary) AS mySum  
FROM Staff  
GROUP BY branchNo  
ORDER BY branchNo;
```

Restricted Groupings – HAVING clause

- HAVING clause is designed for use with GROUP BY to restrict groups that appear in final result table.
- Similar to WHERE, but WHERE filters individual rows whereas HAVING filters groups.
- Column names in HAVING clause must also appear in the GROUP BY list or be contained within an aggregate function.

Select - Use of HAVING

For each branch with more than 1 member of staff, find number of staff in each branch and sum of their salaries.

```
SELECT branchNo,  
        COUNT(staffNo) AS myCount,  
        SUM(salary) AS mySum  
FROM Staff  
GROUP BY branchNo  
HAVING COUNT(staffNo) > 1  
ORDER BY branchNo;
```

Subqueries

- Some SQL statements can have a SELECT embedded within them.
- A subselect can be used in WHERE and HAVING clauses of an outer SELECT, where it is called a *subquery* or *nested query*.
- Subselects may also appear in INSERT, UPDATE, and DELETE statements.

Select - Subquery with Equality

List staff who work in branch at '163 Main St'.

```
SELECT staffNo, fName, lName, position
FROM Staff
WHERE branchNo =
      (SELECT branchNo
       FROM Branch
       WHERE street = '163 Main St');
```

Select - Subquery with Equality

- Inner SELECT finds branch number for branch at '163 Main St' ('B003').
- Outer SELECT then retrieves details of all staff who work at this branch.
- Outer SELECT then becomes:

```
SELECT staffNo, fName, lName, position  
FROM Staff  
WHERE branchNo = 'B003';
```

Select - Subquery with Aggregate

List all staff whose salary is greater than the average salary, and show by how much.

```
SELECT staffNo, fName, lName, position,  
       salary - (SELECT AVG(salary) FROM Staff)  
       As SalDiff  
FROM Staff  
WHERE salary >  
       (SELECT AVG(salary)  
        FROM Staff);
```

Select - Subquery with Aggregate

- Cannot write 'WHERE salary > AVG(salary)'
- Instead, use subquery to find average salary (17000), and then use outer SELECT to find those staff with salary greater than this:

```
SELECT staffNo, fName, lName, position,  
       salary - 17000 As salDiff  
FROM Staff  
WHERE salary > 17000;
```

Subquery Rules

- ORDER BY clause may not be used in a subquery (although it may be used in outermost SELECT).
- Subquery SELECT list must consist of a single column name or expression, except for subqueries that use EXISTS.
- By default, column names refer to table name in FROM clause of subquery. Can refer to a table in FROM using an *alias*.

Subquery Rules

- When subquery is an operand in a comparison, subquery must appear on right-hand side.
- A subquery may not be used as an operand in an expression.

Select - Nested subquery: use of IN

List properties handled by staff at '163 Main St'.

```
SELECT propertyNo, street, city, postcode, type,  
       rooms, rent  
FROM PropertyForRent  
WHERE staffNo IN  
      (SELECT staffNo  
       FROM Staff  
       WHERE branchNo =  
             (SELECT branchNo  
              FROM Branch  
              WHERE street = '163 Main St'));
```

ANY and ALL

- ANY and ALL may be used with subqueries that produce a single column of numbers.
- With ALL, condition will only be true if it is satisfied by *all* values produced by subquery.
- With ANY, condition will be true if it is satisfied by *any* values produced by subquery.
- If subquery is empty, ALL returns true, ANY returns false.
- SOME may be used in place of ANY.

Select - Use of ANY/SOME

Find staff whose salary is larger than salary of at least one member of staff at branch B003.

```
SELECT staffNo, fName, lName, position,  
salary  
FROM Staff  
WHERE salary > SOME  
    (SELECT salary  
     FROM Staff  
     WHERE branchNo = 'B003');
```

Use of ALL

Find staff whose salary is larger than salary of every member of staff at branch B003.

```
SELECT    staffNo,    fName,    lName,  
position, salary  
FROM Staff  
WHERE salary > ALL  
            (SELECT salary  
              FROM Staff  
              WHERE branchNo = 'B003');
```

Functions

- Functions always have the same basic syntax:
 <function name>(function arguments)
- There are hundreds of built-in functions.
- We will be concerned with two broad types of functions:
 - Scalar functions
 - Aggregate functions

Scalar Functions

- Scalar functions operate on a single row at a time.
- Here is a list of scalar functions used in this chapter.

Function Name	Description
GETDATE()	Returns current date and time
MONTH	Returns the month as an integer (1 to 12) from a Date value
YEAR	Returns the Year as a four-digit integer from a date value

Aggregate Functions

- Aggregate functions operate on multiple rows at a time.
- Here is a table of common aggregate functions:

Aggregate Function	Description
COUNT	Counts the number of values : COUNT(*) counts all the rows. COUNT(columnName) counts all the values in the column but ignores nulls
SUM	Sums or totals numeric values: SUM (InStock)
AVG	Returns the mean average of a set of numeric values: AVG(Price). By default nulls are ignored.
MAX	Returns the highest value in a set of numeric or datetime values: MAX(price)
MIN	Returns the smallest value in a set of numeric or datetime values: MIN(Price)

Using Distinct in Aggregate Functions

- You can use the `DISTINCT` keyword with aggregate functions.
- Doing so means the function will ignore duplicate values in its calculation.

```
SELECT COUNT(DISTINCT StudentKey) AS [Unduplicated]  
FROM Session
```


Group By

- When a SELECT clause includes an aggregate function and columns that are not a part of that function, you must use the GROUP BY keywords to group by each of the non-included columns.
- This is necessary because you are mixing functions that operate on multiple rows with columns that refer to values in individual rows only.

Group By Example

```
SELECT Key, COUNT(SessionTimeKey) AS [Total  
Sessions]  
FROM Session  
GROUP BY Key
```

Having

- The HAVING keyword is used when there is an aggregate function in the criteria of a query.

```
SELECT Key, COUNT(SessionTimeKey) AS [Total Sessions]
FROM Session
GROUP BY Key
HAVING COUNT(SessionTimeKey)<4
```

EXISTS and NOT EXISTS

- EXISTS and NOT EXISTS are for use only with subqueries.
- Produce a simple true/false result.
- True if and only if there exists at least one row in result table returned by subquery.
- False if subquery returns an empty result table.
- NOT EXISTS is the opposite of EXISTS.

EXISTS and NOT EXISTS

- As (NOT) EXISTS check only for existence or non-existence of rows in subquery result table, subquery can contain any number of columns.
- Common for subqueries following (NOT) EXISTS to be of form:

(SELECT * ...)

Query using EXISTS

Find all staff who work in a London branch.

```
SELECT staffNo, fName, lName, position
FROM Staff s
WHERE EXISTS
(SELECT *
FROM Branch b
WHERE s.branchNo = b.branchNo AND
      city = 'London');
```

Query using EXISTS

- Note, search condition `s.branchNo = b.branchNo` is necessary to consider correct branch record for each member of staff.
- If omitted, would get all staff records listed out because subquery:

```
SELECT * FROM Branch WHERE city='London'
```

- would always be true and query would be:

```
SELECT staffNo, fName, lName, position FROM Staff  
WHERE true;
```

Query using EXISTS

- Could also write this query using join construct:

```
SELECT    staffNo,    fName,    lName,  
          position  
FROM Staff s, Branch b  
WHERE s.branchNo = b.branchNo AND  
      city = 'London';
```


Joins

- In database design and normalization, the data are broken into several discrete tables.
- Joins are the mechanism for recombining the data into one result set.
- We will look at three kinds of joins:
 - Inner joins
 - Equi joins
 - Outer joins

Multi-Table Queries

- Can use subqueries provided result columns come from same table.
- If result columns come from more than one table must use a join.
- To perform join, include more than one table in FROM clause.
- Use comma as separator and typically include WHERE clause to specify join column(s).

Multi-Table Queries

- Also possible to use an alias for a table named in FROM clause.
- Alias is separated from table name with a space.
- Alias can be used to qualify column names when there is ambiguity.

Simple Join

List names of all clients who have viewed a property along with any comment supplied.

```
SELECT c.clientNo, fName, lName,  
       propertyNo, comment  
FROM Client c, Viewing v  
WHERE c.clientNo = v.clientNo;
```

Simple Join

- Only those rows from both tables that have identical values in the clientNo columns (c.clientNo = v.clientNo) are included in result.
- Equivalent to equi-join in relational algebra.

clientNo	fName	lName	propertyNo	comment
CR56	Aline	Stewart	PG36	too small
CR56	Aline	Stewart	PA14	
CR56	Aline	Stewart	PG4	
CR62	Mary	Tregear	PA14	no dining room
CR76	John	Kay	PG4	too remote

Basic INNER JOIN Syntax

```
SELECT <column1, column2>  
FROM <table1>  
INNER JOIN <table2>  
ON <table1>.<column>=<table2>.<column>
```

Inner Joins

- Inner joins return related records from each of the tables joined.

```
SELECT LastName,  
FirstName,  
SessionDateKey,  
SessionTimeKey,  
StudentKey  
SessionStatus  
FROM  
INNER JOIN Session  
ON .Key = Session.Key
```

Alternative JOIN Constructs

- SQL provides alternative ways to specify joins:

FROM Client c JOIN Viewing v ON c.clientNo = v.clientNo

FROM Client JOIN Viewing USING clientNo

FROM Client NATURAL JOIN Viewing

- In each case, FROM replaces original FROM and WHERE. However, first produces table with two identical clientNo columns.

Equi Joins

- Equi joins present an alternative way to perform inner joins. Some older RDMSs only support this alternative form. The example below also uses an alias for the table name.

```
SELECT t.Key,  
       LastName,  
       FirstName,  
       SessionDateKey,  
       SessionTimeKey,  
       StudentKey  
FROM   t,  
       Session s  
WHERE  t.Key = s.Key  
AND    LastName = 'Brown'
```

Three Table Join

For each branch, list staff who manage properties, including city in which branch is located and properties they manage.

```
SELECT    b.branchNo,    b.city,    s.staffNo,
fName, lName,
           propertyNo
FROM Branch b, Staff s, PropertyForRent p
WHERE b.branchNo = s.branchNo AND
      s.staffNo = p.staffNo
ORDER    BY    b.branchNo,    s.staffNo,
propertyNo;
```

Sorting a join

For each branch, list numbers and names of staff who manage properties, and properties they manage.

```
SELECT s.branchNo, s.staffNo, fName, lName,  
       propertyNo  
FROM Staff s, PropertyForRent p  
WHERE s.staffNo = p.staffNo  
ORDER BY s.branchNo, s.staffNo, propertyNo;
```

Multiple Grouping Columns

Find number of properties handled by each staff member.

```
SELECT s.branchNo, s.staffNo, COUNT(*) AS  
myCount  
FROM Staff s, PropertyForRent p  
WHERE s.staffNo = p.staffNo  
GROUP BY s.branchNo, s.staffNo  
ORDER BY s.branchNo, s.staffNo;
```

Computing a Join

Procedure for generating results of a join are:

1. Form Cartesian product of the tables named in FROM clause.
2. If there is a WHERE clause, apply the search condition to each row of the product table, retaining those rows that satisfy the condition.
3. For each remaining row, determine value of each item in SELECT list to produce a single row in result table.

Computing a Join

4. If DISTINCT has been specified, eliminate any duplicate rows from the result table.
 5. If there is an ORDER BY clause, sort result table as required.
- SQL provides special format of SELECT for Cartesian product:

```
SELECT [DISTINCT | ALL]    {*} | columnList}  
FROM Table1 CROSS JOIN Table2
```

Outer Joins

- If one row of a joined table is unmatched, row is omitted from result table.
- Outer join operations retain rows that do not satisfy the join condition.
- Consider following tables:

branchNo	bCity
B003	Glasgow
B004	Bristol
B002	London

propertyNo	pCity
PA14	Aberdeen
PL94	London
PG4	Glasgow

Outer Joins

- The (inner) join of these two tables:

```
SELECT b.*, p.*  
FROM Branch1 b, PropertyForRent1 p  
WHERE b.bCity = p.pCity;
```

branchNo	bCity	propertyNo	pCity
B003	Glasgow	PG4	Glasgow
B002	London	PL94	London

Outer Joins

- Result table has two rows where cities are same.
- There are no rows corresponding to branches in Bristol and Aberdeen.
- To include unmatched rows in result table, use an Outer join.

Left Outer Join

List branches and properties that are in same city along with **any unmatched properties**.

```
SELECT b.*, p.*  
FROM Branch1 b LEFT JOIN  
    PropertyForRent1 p ON b.bCity =  
    p.pCity;
```

Left Outer Join

- Includes those rows of first (left) table unmatched with rows from second (right) table.
- Columns from second table are filled with NULLs.

branchNo	bCity	propertyNo	pCity
B003	Glasgow	PG4	Glasgow
B004	Bristol	NULL	NULL
B002	London	PL94	London

Right Outer Join

List branches and properties in same city and any unmatched properties.

```
SELECT b.*, p.*  
FROM Branch1 b RIGHT JOIN  
PropertyForRent1 p ON b.bCity =  
p.pCity;
```

Right Outer Join

- Right Outer join includes those rows of second (right) table that are unmatched with rows from first (left) table.
- Columns from first table are filled with NULLs.

branchNo	bCity	propertyNo	pCity
NULL	NULL	PA14	Aberdeen
B003	Glasgow	PG4	Glasgow
B002	London	PL94	London

Full Outer Join

List branches and properties in same city and **any unmatched branches or properties. (on both sides)**

```
SELECT b.*, p.*  
FROM Branch1 b FULL JOIN  
PropertyForRent1 p ON b.bCity = p.pCity;
```

Full Outer Join

- Includes rows that are unmatched in both tables.
- Unmatched columns are filled with NULLs.

branchNo	bCity	propertyNo	pCity
NULL	NULL	PA14	Aberdeen
B003	Glasgow	PG4	Glasgow
B004	Bristol	NULL	NULL
B002	London	PL94	London

OUTER JOIN Syntax

Outer joins return records that are not matched. The following query returns s that have no sessions scheduled.

```
SELECT <column1>, <column2>  
FROM <table1>  
LEFT OUTER JOIN <table2>  
ON <table1>.<column>=<table2>.<column>
```


Outer Join Example

```
SELECT t.Key,  
       LastName,  
       SessionDateKey  
FROM   t  
LEFT OUTER JOIN Session s  
ON t.Key = s.Key  
WHERE SessionDateKey IS Null
```

Inserts

To insert a record into a table, you use the following syntax:

```
INSERT INTO <tablename>(<ColumnName>,  
<columnName>, ...)  
VALUES(<value1>, <value2>, ...)
```

Updates

Updates allow you to change existing records. The syntax is:

```
UPDATE <TableName>  
SET <ColumnName> = <New Value>,  
  <ColumnName>=<new value>  
WHERE <ColumnName> = <criteria>
```

Deletes

- Deletes allow you to remove a record from a table:

```
DELETE FROM <TableName>  
WHERE <columnName> = <criteria>
```

Deletes and Updates

- Deletes and updates are dangerous. If you do not specify a criteria, the update or delete will be applied to all the rows in a table.
- Also, referential integrity may prevent a deletion. You cannot delete a parent that has children in another table.

SubQuery Example

```
SELECT DISTINCT COUNT(*) AS Total,  
(SELECT COUNT(*)  
FROM Session  
WHERE SessionStatus='NS') AS NoShow,  
(SELECT COUNT(*)  
FROM Session  
WHERE SessionStatus='c') AS Completed  
FROM Session
```

This example shows subqueries used in the SELECT clause to return Aggregate values.

Locating Duplicates

```
SELECT Lastname, firstname, email, phone,  
COUNT(*) AS [duplicates]  
FROM contact  
GROUP BY Lastname, firstName, email, Phone  
HAVING COUNT(*) >1
```

This SQL finds duplicate values in in a table.

Documentation: Testing Plans

- When testing the database, you should document all your SQL queries and their results.
- On the next slide is a sample of a test table, showing the test and results.

Union, Intersect, and Difference

- Can use normal set operations of Union, Intersection, and Difference to combine results of two or more queries into a single result table.
- Union of two tables, A and B, is table containing all rows in either A or B or both.
- Intersection is table containing all rows common to both A and B.
- Difference is table containing all rows in A but not in B.
- Two tables must be *union compatible*.


Union, Intersect, and Difference

- Format of set operator clause in each case is:

op [ALL] [CORRESPONDING [BY {column1 [, ...]}]]

- If CORRESPONDING BY specified, set operation performed on the named column(s).
- If CORRESPONDING specified but not BY clause, operation performed on common columns.
- If ALL specified, result can include duplicate rows.

Union, Intersect, and Difference



Use of UNION

List all cities where there is either a branch office or a property.

```
(SELECT city
FROM Branch
WHERE city IS NOT NULL) UNION
(SELECT city
FROM PropertyForRent
WHERE city IS NOT NULL);
```

Use of UNION

- Or

```
(SELECT *  
FROM Branch  
WHERE city IS NOT NULL)  
UNION CORRESPONDING BY city  
(SELECT *  
FROM PropertyForRent  
WHERE city IS NOT NULL);
```

Union Example

```
INSERT INTO Contact(LastName, FirstName, Email, Phone)
SELECT StudentLastName AS LastName,
StudentFirstName AS FirstName,
StudentEmail AS Email,
StudentPhone AS Phone
FROM Student
WHERE StudentEmail IS NOT NULL
UNION
SELECT LastName,
FirstName,
Email,
Phone
FROM
WHERE Email IS NOT NULL
```

This UNION query joins the tables Student and into a single result and writes them to the table Contact.

Use of UNION

- Produces result tables from both queries and merges both tables together.

city
London
Glasgow
Aberdeen
Bristol

Use of INTERSECT

List all cities where there is both a branch office and a property.

```
(SELECT city FROM Branch)  
INTERSECT  
(SELECT city FROM PropertyForRent);
```


Use of INTERSECT

- Or

```
(SELECT * FROM Branch)  
INTERSECT CORRESPONDING BY city  
(SELECT * FROM PropertyForRent);
```

city
Aberdeen
Glasgow
London

Use of INTERSECT

- Could rewrite this query without INTERSECT operator:

```
SELECT b.city  
FROM Branch b PropertyForRent p  
WHERE b.city = p.city;
```

- Or:

```
SELECT DISTINCT city FROM Branch b  
WHERE EXISTS  
(SELECT * FROM PropertyForRent p  
WHERE p.city = b.city);
```

Use of EXCEPT

List of all cities where there is a branch office but no properties.

```
(SELECT city FROM Branch)
EXCEPT
(SELECT city FROM PropertyForRent);
```

- Or

```
(SELECT * FROM Branch)
EXCEPT CORRESPONDING BY city
(SELECT * FROM PropertyForRent);
```

city
Bristol

Use of EXCEPT

- Could rewrite this query without EXCEPT:

```
SELECT DISTINCT city FROM Branch
WHERE city NOT IN
  (SELECT city FROM PropertyForRent);
```

- Or

```
SELECT DISTINCT city FROM Branch b
WHERE NOT EXISTS
  (SELECT * FROM PropertyForRent p
   WHERE p.city = b.city);
```

INSERT

```
INSERT INTO TableName [ (columnList) ]  
VALUES (dataValueList)
```

- *columnList* is optional; if omitted, SQL assumes a list of all columns in their original CREATE TABLE order.
- Any columns omitted must have been declared as NULL when table was created, unless DEFAULT was specified when creating column.

INSERT

- *dataValueList* must match *columnList* as follows:
 - number of items in each list must be same;
 - must be direct correspondence in position of items in two lists;
 - data type of each item in *dataValueList* must be compatible with data type of corresponding column.

INSERT ... VALUES

Insert a new row into Staff table supplying data for all columns.

```
INSERT INTO Staff
VALUES      ('SG16',      'Alan',      'Brown',
            'Assistant', 'M', Date '1957-05-25', 8300,
            'B003');
```

INSERT using Defaults

Insert a new row into Staff table supplying data for all mandatory columns.

```
INSERT INTO Staff (staffNo, fName, lName,  
                  position,  
                  salary, branchNo)  
VALUES ('SG44', 'Anne', 'Jones',  
        'Assistant', 8100, 'B003');
```

- Or

```
INSERT INTO Staff  
VALUES ('SG44', 'Anne', 'Jones', 'Assistant',  
        NULL,  
        NULL, 8100, 'B003');
```


INSERT ... SELECT

- Second form of INSERT allows multiple rows to be copied from one or more tables to another:

```
INSERT INTO TableName [ (columnList) ]  
    SELECT ...
```

INSERT ... SELECT

Assume there is a table StaffPropCount that contains names of staff and number of properties they manage:

```
StaffPropCount(staffNo,      fName,      lName,  
               propCnt)
```

Populate StaffPropCount using Staff and
PropertyForRent tables.

INSERT ... SELECT

```
INSERT INTO StaffPropCount
  (SELECT s.staffNo, fName, lName, COUNT(*)
   FROM Staff s, PropertyForRent p
   WHERE s.staffNo = p.staffNo
   GROUP BY s.staffNo, fName, lName)
UNION
  (SELECT staffNo, fName, lName, 0
   FROM Staff
   WHERE staffNo NOT IN
     (SELECT DISTINCT staffNo
      FROM PropertyForRent));
```

UPDATE

```
UPDATE TableName  
SET columnName1 = dataValue1  
    [, columnName2 = dataValue2...]  
[WHERE searchCondition]
```

- *TableName* can be name of a base table or an updatable view.
- SET clause specifies names of one or more columns that are to be updated.

UPDATE

- WHERE clause is optional:
 - if omitted, named columns are updated for all rows in table;
 - if specified, only those rows that satisfy *searchCondition* are updated.
- New *dataValue(s)* must be compatible with data type for corresponding column.

UPDATE All Rows

Give all staff a 3% pay increase.

```
UPDATE Staff  
SET salary = salary*1.03;
```

Give all Managers a 5% pay increase.

```
UPDATE Staff  
SET salary = salary*1.05  
WHERE position = 'Manager';
```

UPDATE Multiple Columns

Promote David Ford (staffNo='SG14') to Manager and change his salary to £18,000.

```
UPDATE Staff
SET position = 'Manager', salary = 18000
WHERE staffNo = 'SG14';
```

DELETE

```
DELETE FROM TableName  
[WHERE searchCondition]
```

- *TableName* can be name of a base table or an updatable view.
- *searchCondition* is optional; if omitted, all rows are deleted from table. This does not delete table. If *search_condition* is specified, only those rows that satisfy condition are deleted.

DELETE Specific Rows

Delete all viewings that relate to property PG4.

```
DELETE FROM Viewing  
WHERE propertyNo = 'PG4';
```

Delete all records from the Viewing table.

```
DELETE FROM Viewing;
```

ISO SQL Data Types

DATA TYPE	DECLARATIONS				
boolean	BOOLEAN				
character	CHAR	VARCHAR			
bit [†]	BIT	BIT VARYING			
exact numeric	NUMERIC	DECIMAL	INTEGER	SMALLINT	BIGINT
approximate numeric	FLOAT	REAL	DOUBLE PRECISION		
datetime	DATE	TIME	TIMESTAMP		
interval	INTERVAL				
large objects	CHARACTER LARGE OBJECT		BINARY LARGE OBJECT		

[†]BIT and BIT VARYING have been removed from the SQL:2003 standard.

Data Definition

- SQL DDL allows database objects such as schemas, domains, tables, views, and indexes to be created and destroyed.
- Main SQL DDL statements are:

CREATE SCHEMA	DROP SCHEMA
CREATE/ALTER DOMAIN	DROP DOMAIN
CREATE/ALTER TABLE	DROP TABLE
CREATE VIEW	DROP VIEW

- Many DBMSs also provide:

CREATE INDEX	DROP INDEX
--------------	------------

Data Definition

- Relations and other database objects exist in an *environment*.
- Each environment contains one or more *catalogs*, and each catalog consists of set of schemas.
- Schema is named collection of related database objects.
- Objects in a schema can be tables, views, domains, assertions, collations, translations, and character sets. All have same owner.

CREATE SCHEMA

```
CREATE SCHEMA [Name |  
    AUTHORIZATION CreatorId ]  
DROP SCHEMA Name [RESTRICT | CASCADE ]
```

- With RESTRICT (default), schema must be empty or operation fails.
- With CASCADE, operation cascades to drop all objects associated with schema in order defined above. If any of these operations fail, DROP SCHEMA fails.

CREATE TABLE

```
CREATE TABLE TableName
{(colName dataType [NOT NULL] [UNIQUE]
[DEFAULT defaultOption]
[CHECK searchCondition] [,...]}
[PRIMARY KEY (listOfColumns),]
{[UNIQUE (listOfColumns),] [...,]}
{[FOREIGN KEY (listOfFKColumns)
REFERENCES                               ParentTableName
[(listOfCKColumns)],
[ON UPDATE referentialAction]
[ON DELETE referentialAction ]] [,...]}
{[CHECK (searchCondition)] [,...]} )
```

CREATE TABLE

- Creates a table with one or more columns of the specified *dataType*.
- With NOT NULL, system rejects any attempt to insert a null in the column.
- Can specify a DEFAULT value for the column.
- Primary keys should always be specified as NOT NULL.
- FOREIGN KEY clause specifies FK along with the referential action.

CREATE TABLE

```
CREATE TABLE PropertyForRent (  
    propertyNo PNumber NOT NULL, ...  
    rooms      PRooms  NOT NULL  DEFAULT 4,  
    rent       PRent    NOT NULL,  DEFAULT 600,  
    ownerNo    OwnerNumber NOT NULL,  
    staffNo    StaffNumber  
                Constraint StaffNotHandlingTooMuch ...  
    branchNo   BranchNumber NOT NULL,  
    PRIMARY KEY (propertyNo),  
    FOREIGN KEY (staffNo) REFERENCES Staff  
                ON DELETE SET NULL ON UPDATE CASCADE ....);
```


ALTER TABLE

- Add a new column to a table.
- Drop a column from a table.
- Add a new table constraint.
- Drop a table constraint.
- Set a default for a column.
- Drop a default for a column.

ALTER TABLE

Change Staff table by removing default of 'Assistant' for position column and setting default for sex column to female ('F').

```
ALTER TABLE Staff  
ALTER position DROP DEFAULT;  
ALTER TABLE Staff  
ALTER sex SET DEFAULT 'F';
```

ALTER TABLE

Remove constraint from PropertyForRent that staff are not allowed to handle more than 100 properties at a time. Add new column to Client table.

```
ALTER TABLE PropertyForRent
    DROP CONSTRAINT StaffNotHandlingTooMuch;

ALTER TABLE Client
    ADD prefNoRooms PRooms;
```

DROP TABLE

`DROP TABLE TableName [RESTRICT | CASCADE]`

e.g. `DROP TABLE PropertyForRent;`

- Removes named table and all rows within it.
- With RESTRICT, if any other objects depend for their existence on continued existence of this table, SQL does not allow request.
- With CASCADE, SQL drops all dependent objects (and objects dependent on these objects).

Views

View

Dynamic result of one or more relational operations operating on base relations to produce another relation.

- Virtual relation that does not necessarily actually exist in the database but is produced upon request, at time of request.

Views

- Contents of a view are defined as a query on one or more base relations.
- With view resolution, any operations on view are automatically translated into operations on relations from which it is derived.
- With view materialization, the view is stored as a temporary table, which is maintained as the underlying base tables are updated.

SQL - CREATE VIEW

```
CREATE VIEW ViewName [ (newColumnName [,...]) ]  
    AS subselect  
    [WITH [CASCADED | LOCAL] CHECK OPTION]
```

- Can assign a name to each column in view.
- If list of column names is specified, it must have same number of items as number of columns produced by *subselect*.
- If omitted, each column takes name of corresponding column in *subselect*.

SQL - CREATE VIEW

- List must be specified if there is any ambiguity in a column name.
- The *subselect* is known as the defining query.
- WITH CHECK OPTION ensures that if a row fails to satisfy WHERE clause of defining query, it is not added to underlying base table.
- Need SELECT privilege on all tables referenced in subselect and USAGE privilege on any domains used in referenced columns.

Create Horizontal View

Create view so that manager at branch B003 can only see details for staff who work in his or her office.

```
CREATE VIEW Manager3Staff  
AS      SELECT *  
        FROM Staff  
        WHERE branchNo = 'B003';
```

staffNo	fName	lName	position	sex	DOB	salary	branchNo
SG37	Ann	Beech	Assistant	F	10-Nov-60	12000.00	B003
SG14	David	Ford	Supervisor	M	24-Mar-58	18000.00	B003
SG5	Susan	Brand	Manager	F	3-Jun-40	24000.00	B003

Create Vertical View

Create view of staff details at branch B003 excluding salaries.

```
CREATE VIEW Staff3  
AS SELECT staffNo, fName, lName,  
position, sex  
FROM Staff  
WHERE branchNo = 'B003';
```

staffNo	fName	lName	position	sex
SG37	Ann	Beech	Assistant	F
SG14	David	Ford	Supervisor	M
SG5	Susan	Brand	Manager	F

Grouped and Joined Views

Create view of staff who manage properties for rent, including branch number they work at, staff number, and number of properties they manage.

```
CREATE VIEW StaffPropCnt (branchNo,  
staffNo, cnt)  
AS SELECT s.branchNo, s.staffNo, COUNT(*)  
FROM Staff s, PropertyForRent p  
WHERE s.staffNo = p.staffNo  
GROUP BY s.branchNo, s.staffNo;
```

SQL - DROP VIEW

`DROP VIEW ViewName [RESTRICT | CASCADE]`

- Causes definition of view to be deleted from database.
- For example:

`DROP VIEW Manager3Staff;`

SQL - DROP VIEW

- With CASCADE, all related dependent objects are deleted; i.e. any views defined on view being dropped.
- With RESTRICT (default), if any other objects depend for their existence on continued existence of view being dropped, command is rejected.

View Resolution

Count number of properties managed by each member at branch B003.

```
SELECT staffNo, cnt  
FROM StaffPropCnt  
WHERE branchNo = 'B003'  
ORDER BY staffNo;
```

View Resolution

- (a) View column names in SELECT list are translated into their corresponding column names in the defining query:

```
SELECT s.staffNo As staffNo, COUNT(*) As cnt
```

- (b) View names in FROM are replaced with corresponding FROM lists of defining query:

```
FROM Staff s, PropertyForRent p
```

View Resolution

- (c) WHERE from user query is combined with WHERE of defining query using AND:

WHERE s.staffNo = p.staffNo AND branchNo = 'B003'

- (d) GROUP BY and HAVING clauses copied from defining query:

GROUP BY s.branchNo, s.staffNo

- (e) ORDER BY copied from query with view column name translated into defining query column name

ORDER BY s.staffNo

View Resolution

(f) Final merged query is now executed to produce the result:

```
SELECT s.staffNo AS staffNo, COUNT(*)  
AS cnt  
FROM Staff s, PropertyForRent p  
WHERE s.staffNo = p.staffNo AND  
branchNo = 'B003'  
GROUP BY s.branchNo, s.staffNo  
ORDER BY s.staffNo;
```

Restrictions on Views

SQL imposes several restrictions on creation and use of views.

(a) If column in view is based on an aggregate function:

- Column may appear only in SELECT and ORDER BY clauses of queries that access view.
- Column may not be used in WHERE nor be an argument to an aggregate function in any query based on view.

Restrictions on Views

- For example, following query would fail:

```
SELECT COUNT(cnt)
FROM StaffPropCnt;
```

- Similarly, following query would also fail:

```
SELECT *
FROM StaffPropCnt
WHERE cnt > 2;
```

Restrictions on Views

- (b) Grouped view may never be joined with a base table or a view.
- For example, StaffPropCnt view is a grouped view, so any attempt to join this view with another table or view fails.

View Updatability

- All updates to base table reflected in all views that encompass base table.
- Similarly, may expect that if view is updated then base table(s) will reflect change.

View Updatability

- However, consider again view StaffPropCnt.
- If we tried to insert record showing that at branch B003, SG5 manages 2 properties:

```
INSERT INTO StaffPropCnt  
VALUES ('B003', 'SG5', 2);
```

- Have to insert 2 records into PropertyForRent showing which properties SG5 manages. However, do not know which properties they are; i.e. do not know primary keys!

View Updatability

- If change definition of view and replace count with actual property numbers:

```
CREATE VIEW StaffPropList (branchNo,  
                           staffNo, propertyNo)  
AS SELECT s.branchNo, s.staffNo, p.propertyNo  
   FROM Staff s, PropertyForRent p  
  WHERE s.staffNo = p.staffNo;
```

View Updatability

- Now try to insert the record:

```
INSERT INTO StaffPropList  
VALUES ('B003', 'SG5', 'PG19');
```

- Still problem, because in PropertyForRent all columns except postcode/staffNo are not allowed nulls.
- However, have no way of giving remaining non-null columns values.

View Updatability

- ISO specifies that a view is updatable if and only if:
 - DISTINCT is not specified.
 - Every element in SELECT list of defining query is a column name and no column appears more than once.
 - FROM clause specifies only one table, excluding any views based on a join, union, intersection or difference.
 - No nested SELECT referencing outer table.
 - No GROUP BY or HAVING clause.
 - Also, every row added through view must not violate integrity constraints of base table.

Updatable View

For view to be updatable, DBMS must be able to trace any row or column back to its row or column in the source table.

WITH CHECK OPTION

- Rows exist in a view because they satisfy WHERE condition of defining query.
- If a row changes and no longer satisfies condition, it disappears from the view.
- New rows appear within view when insert/update on view cause them to satisfy WHERE condition.
- Rows that enter or leave a view are called *migrating rows*.
- WITH CHECK OPTION prohibits a row migrating out of the view.

WITH CHECK OPTION

- LOCAL/CASCADED apply to view hierarchies.
- With LOCAL, any row insert/update on view and any view directly or indirectly defined on this view must not cause row to disappear from view unless row also disappears from derived view/table.
- With CASCADED (default), any row insert/ update on this view and on any view directly or indirectly defined on this view must not cause row to disappear from the view.

WITH CHECK OPTION

```
CREATE VIEW Manager3Staff  
AS      SELECT *  
      FROM Staff  
      WHERE branchNo = 'B003'  
WITH CHECK OPTION;
```

- Cannot update branch number of row B003 to B002 as this would cause row to migrate from view.
- Also cannot insert a row into view with a branch number that does not equal B003.

WITH CHECK OPTION

- Now consider the following:

```
CREATE VIEW LowSalary
  AS SELECT * FROM Staff WHERE salary > 9000;

CREATE VIEW HighSalary
  AS SELECT * FROM LowSalary
    WHERE salary > 10000
  WITH LOCAL CHECK OPTION;

CREATE VIEW Manager3Staff
  AS SELECT * FROM HighSalary
    WHERE branchNo = 'B003';
```

WITH CHECK OPTION

```
UPDATE Manager3Staff  
SET salary = 9500  
WHERE staffNo = 'SG37';
```

- This update would fail: although update would cause row to disappear from HighSalary, row would not disappear from LowSalary.
- However, if update tried to set salary to 8000, update would succeed as row would no longer be part of LowSalary.

WITH CHECK OPTION

- If HighSalary had specified WITH CASCADED CHECK OPTION, setting salary to 9500 or 8000 would be rejected because row would disappear from HighSalary.
- To prevent anomalies like this, each view should be created using WITH CASCADED CHECK OPTION.

Advantages of Views

- Data independence
- Currency
- Improved security
- Reduced complexity
- Convenience
- Customization
- Data integrity

Disadvantages of Views

- Update restriction
- Structure restriction
- Performance

View Materialization

- View resolution mechanism may be slow, particularly if view is accessed frequently.
- View materialization stores view as temporary table when view is first queried.
- Thereafter, queries based on materialized view can be faster than recomputing view each time.
- Difficulty is maintaining the currency of view while base tables(s) are being updated.

View Maintenance

- View maintenance aims to apply only those changes necessary to keep view current.
- Consider following view:

```
CREATE VIEW StaffPropRent(staffNo)
AS SELECT DISTINCT staffNo
   FROM PropertyForRent
   WHERE branchNo = 'B003' AND
         rent > 400;
```

View Materialization

- If insert row into PropertyForRent with rent ≤ 400 then view would be unchanged.
- If insert row for property PG24 at branch B003 with staffNo = SG19 and rent = 550, then row would appear in materialized view.
- If insert row for property PG54 at branch B003 with staffNo = SG37 and rent = 450, then no new row would need to be added to materialized view.
- If delete property PG24, row should be deleted from materialized view.
- If delete property PG54, then row for PG37 should not be deleted (because of existing property PG21).

Transactions

- SQL defines transaction model based on COMMIT and ROLLBACK.
- Transaction is logical unit of work with one or more SQL statements guaranteed to be atomic with respect to recovery.
- An SQL transaction automatically begins with a *transaction-initiating* SQL statement (e.g., SELECT, INSERT).
- Changes made by transaction are not visible to other concurrently executing transactions until transaction completes.

Transactions

- Transaction can complete in one of four ways:
 - COMMIT ends transaction successfully, making changes permanent.
 - ROLLBACK aborts transaction, backing out any changes made by transaction.
 - For programmatic SQL, successful program termination ends final transaction successfully, even if COMMIT has not been executed.
 - For programmatic SQL, abnormal program end aborts transaction.

Transactions

- New transaction starts with next transaction-initiating statement.
- SQL transactions cannot be nested.
- SET TRANSACTION configures transaction:

SET TRANSACTION

[READ ONLY | READ WRITE] |

[ISOLATION LEVEL READ UNCOMMITTED |

READ COMMITTED | REPEATABLE READ | SERIALIZABLE]

Immediate and Deferred Integrity Constraints

- Do not always want constraints to be checked immediately, but instead at transaction commit.
- Constraint may be defined as `INITIALLY IMMEDIATE` or `INITIALLY DEFERRED`, indicating mode the constraint assumes at start of each transaction.
- In former case, also possible to specify whether mode can be changed subsequently using qualifier `[NOT] DEFERRABLE`.
- Default mode is `INITIALLY IMMEDIATE`.

Immediate and Deferred Integrity Constraints

- SET CONSTRAINTS statement used to set mode for specified constraints for current transaction:

SET CONSTRAINTS

```
{ALL | constraintName [, . . . ]}  
{DEFERRED | IMMEDIATE}
```

Access Control - Authorization and Ownership

- Authorization identifier is normal SQL identifier used to establish identity of a user. Usually has an associated password.
- Used to determine which objects user may reference and what operations may be performed on those objects.
- Each object created in SQL has an owner, as defined in AUTHORIZATION clause of schema to which object belongs.
- Owner is only person who may know about it.

Privileges

- Actions user permitted to carry out on given base table or view:

SELECT Retrieve data from a table.

INSERT Insert new rows into a table.

UPDATE Modify rows of data in a table.

DELETE Delete rows of data from a table.

REFERENCES Reference columns of named table in integrity constraints.

USAGE Use domains, collations, character sets, and translations.

Privileges

- Can restrict INSERT/UPDATE/REFERENCES to named columns.
- Owner of table must grant other users the necessary privileges using GRANT statement.
- To create view, user must have SELECT privilege on all tables that make up view and REFERENCES privilege on the named columns.

GRANT

```
GRANT {PrivilegeList | ALL PRIVILEGES}  
ON ObjectName  
TO {AuthorizationIdList | PUBLIC}  
[WITH GRANT OPTION]
```

- *PrivilegeList* consists of one or more of above privileges separated by commas.
- ALL PRIVILEGES grants all privileges to a user.

GRANT

- PUBLIC allows access to be granted to all present and future authorized users.
- *ObjectName* can be a base table, view, domain, character set, collation or translation.
- WITH GRANT OPTION allows privileges to be passed on.

Example GRANT

Give Manager full privileges to Staff table.

```
GRANT ALL PRIVILEGES  
ON Staff  
TO Manager WITH GRANT OPTION;
```

Give users Personnel and Director SELECT and UPDATE on column salary of Staff.

```
GRANT SELECT, UPDATE (salary)  
ON Staff  
TO Personnel, Director;
```


GRANT Specific Privileges to PUBLIC

Give all users SELECT on Branch table.

```
GRANT SELECT  
ON Branch  
TO PUBLIC;
```

REVOKE

- REVOKE takes away privileges granted with GRANT.

```
REVOKE [GRANT OPTION FOR]
    {PrivilegeList | ALL PRIVILEGES}
ON ObjectName
FROM {AuthorizationIdList | PUBLIC}
    [RESTRICT | CASCADE]
```

- ALL PRIVILEGES refers to all privileges granted to a user by user revoking privileges.

REVOKE

- GRANT OPTION FOR allows privileges passed on via WITH GRANT OPTION of GRANT to be revoked separately from the privileges themselves.
- REVOKE fails if it results in an abandoned object, such as a view, unless the CASCADE keyword has been specified.
- Privileges granted to this user by other users are not affected.

REVOKE Specific Privileges

Revoke privilege SELECT on Branch table from all users.

```
REVOKE SELECT  
ON Branch  
FROM PUBLIC;
```

Revoke all privileges given to Director on Staff table.

```
REVOKE ALL PRIVILEGES  
ON Staff  
FROM Director;
```

Creating a Trigger

- Triggers are programs that are triggered by an event, typically INSERT, UPDATE, or DELETE.
- They can be used to enforce business rules that referential integrity and constraints alone cannot enforce.
- The basic syntax for creating a trigger is:

```
CREATE TRIGGER <trigger_name> ON <table_name>
[FOR, AFTER, INSTEAD OF] [INSERT, UPDATE,
DELETE]
AS
{SQL Code}
```

Advanced SQL

- SQL is a powerful language and there is much more that can be done with it.
- Subqueries allow a user to embed whole independent SELECT statements in the SELECT clause or as a criterion in the WHERE clause.
- Unions allow a user to blend the results of a two-result set into a single tabular output.
- You can use SQL to find and remove duplicates.
- Indexes help a database administrator speed up query results and optimize the database.

Integrity Enhancement Feature

- Consider five types of integrity constraints:
 - required data
 - domain constraints
 - entity integrity
 - referential integrity
 - general constraints.

Integrity Enhancement Feature

Required Data

position	VARCHAR(10)	NOT NULL
----------	-------------	----------

Domain Constraints

(a) CHECK

sex	CHAR NOT NULL
	CHECK (sex IN ('M', 'F'))

Integrity Enhancement Feature

(b) CREATE DOMAIN

```
CREATE DOMAIN DomainName [AS] dataType  
[DEFAULT defaultOption]  
[CHECK (searchCondition)]
```

For example:

```
CREATE DOMAIN SexType AS CHAR  
CHECK (VALUE IN ('M', 'F'));  
sex SexType NOT NULL
```

Integrity Enhancement Feature

- *searchCondition* can involve a table lookup:

```
CREATE DOMAIN BranchNo AS CHAR(4)  
CHECK (VALUE IN (SELECT branchNo  
                  FROM Branch));
```

- Domains can be removed using DROP DOMAIN:

```
DROP DOMAIN DomainName  
[RESTRICT | CASCADE]
```

IEF - Entity Integrity

- Primary key of a table must contain a unique, non-null value for each row.
- ISO standard supports FOREIGN KEY clause in CREATE and ALTER TABLE statements:

PRIMARY KEY(staffNo)

PRIMARY KEY(clientNo, propertyNo)

- Can only have one PRIMARY KEY clause per table. Can still ensure uniqueness for alternate keys using UNIQUE:

UNIQUE(telNo)

IEF - Referential Integrity

- FK is column or set of columns that links each row in child table containing foreign FK to row of parent table containing matching PK.
- Referential integrity means that, if FK contains a value, that value must refer to existing row in parent table.
- ISO standard supports definition of FKs with FOREIGN KEY clause in CREATE and ALTER TABLE:

FOREIGN KEY(branchNo) REFERENCES Branch

IEF - Referential Integrity

- Any INSERT/UPDATE attempting to create FK value in child table without matching CK value in parent is rejected.
- Action taken attempting to update/delete a CK value in parent table with matching rows in child is dependent on referential action specified using ON UPDATE and ON DELETE subclauses:
 - CASCADE
 - SET NULL
 - SET DEFAULT
 - NO ACTION

IEF - Referential Integrity

CASCADE: Delete row from parent and delete matching rows in child, and so on in cascading manner.

SET NULL: Delete row from parent and set FK column(s) in child to NULL. Only valid if FK columns are NOT NULL.

SET DEFAULT: Delete row from parent and set each component of FK in child to specified default. Only valid if DEFAULT specified for FK columns.

NO ACTION: Reject delete from parent. Default.

IEF - Referential Integrity

FOREIGN KEY (staffNo) REFERENCES Staff
ON DELETE SET NULL

FOREIGN KEY (ownerNo) REFERENCES Owner
ON UPDATE CASCADE

IEF - General Constraints

- Could use CHECK/UNIQUE in CREATE and ALTER TABLE.
- Similar to the CHECK clause, also have:

```
CREATE ASSERTION AssertionName  
CHECK (searchCondition)
```


IEF - General Constraints

```
CREATE ASSERTION StaffNotHandlingTooMuch
CHECK (NOT EXISTS      (SELECT staffNo
                        FROM PropertyForRent
                        GROUP BY staffNo
                        HAVING COUNT(*) > 100))
```